

# (my) Principles for Developing Machine Learning Projects

Hugo Sousa

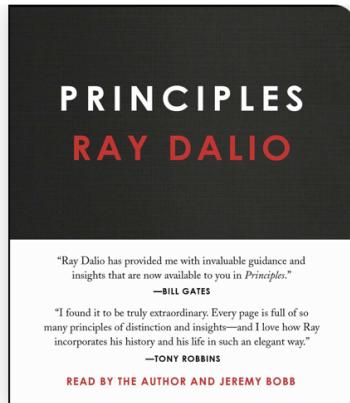
NLP Researcher  
Applied Scientist @Amazon

# Career context

- Bachelor in Physics
  - Master in Applied Mathematics
  - (almost) PhD in Computer Science
- 
- Visiting Student @CMU
  - Internship @Amazon in Palo Alto
- 
- Doing research in NLP and IR

Website: <https://hugosousa.net/>

# (my) Principles for Developing Machine Learning Projects

[Want to Read](#)[Buy on Amazon](#)[Rate this book](#)

# Principles: Life and Work

Ray Dalio (Author/Narrator) 9, Jeremy Bobb (Reader)

4.10 66,619 ratings · 3,833 reviews

Ray Dalio, one of the world's most successful investors and entrepreneurs, shares the unconventional principles that he's developed, refined, and used over the past forty years to create unique results in both life and business—and which any person or organization can adopt to help achieve their goals.

[Show more ▾](#)[Genres](#)[Business](#)[Nonfiction](#)[Self Help](#)[Finance](#)[Leadership](#)[Personal Development](#)[Philosophy](#)[...more](#)

16 pages, Audio CD

First published September 19, 2017

[Book details & editions ▾](#)

# The Principals

- Be obsessed with your data
- Be maniacal about your code
- Constrain as little as possible
- Just do it
- Be bold, be simple

*Be obsessed with your data*

# **GIGO**: “*Garbage in, garbage out*”

- Coined circa **1957** from US mathematicians
- **It doesn't matter how sophisticate your learning mechanism is or how much compute you have if your data is of poor quality**

# The infamous tank story

## Neural Network Follies

*by Neil Fraser, September 1998*

In the 1980s, the Pentagon wanted to harness computer technology to make their tanks harder to attack.

<https://neil.fraser.name/writing/tank/>

- Neil Fraser later admitted that the story "might be apocryphal."
- But a very similar one the Wolf/Husky has been discussed on the “Why should I trust you paper?”

<https://dl.acm.org/doi/10.1145/2939672.2939778>

# The “it” in AI models is the dataset.

Posted on June 10, 2023 by jbetker

I've been at OpenAI for almost a year now. In that time, I've trained a **lot** of generative models. More than anyone really has any right to train. As I've spent these hours observing the effects of tweaking various model configurations and hyperparameters, one thing that has struck me is the similarities in between all the training runs.

It's becoming awfully clear to me that these models are truly approximating their datasets to an incredible degree. What that means is not only that they learn what it means to be a dog or a cat, but the interstitial frequencies between distributions that don't matter, like what photos humans are likely to take or words humans commonly write down.

What this manifests as is – trained on the same dataset for long enough, pretty much every model with enough weights and training time converges to the same point. Sufficiently large diffusion conv-unets produce the same images as ViT generators. AR sampling produces the same images as diffusion.

This is a surprising observation! It implies that model behavior is not determined by architecture, hyperparameters, or optimizer choices. It's determined by your dataset, nothing else. Everything else is a means to an end in efficiently delivery compute to approximating that dataset.

Then, when you refer to "Lambda", "ChatGPT", "Bard", or "Claude" then, it's not the model weights that you are referring to. It's the dataset.

<https://nonint.com/2023/06/10/the-it-in-ai-models-is-the-dataset/>

# Prompt: “Write a metaphor about time”

ChatGPT ▾

Time is a quiet river carving mountains—so gentle you hardly feel it passing, yet so persistent that nothing you know remains unchanged.



☰ Gemini

⌚ Time as a River

Time is a **river**, swift and relentless, flowing ever onward to the sea of the future. We can stand by its banks, watch the moments rush past, but we can never step into the same current twice.



Time is a river that carries us forward on its current—we can dip our hands in to feel each moment pass through our fingers, but we can never hold the water still or swim back upstream to where we've already been.

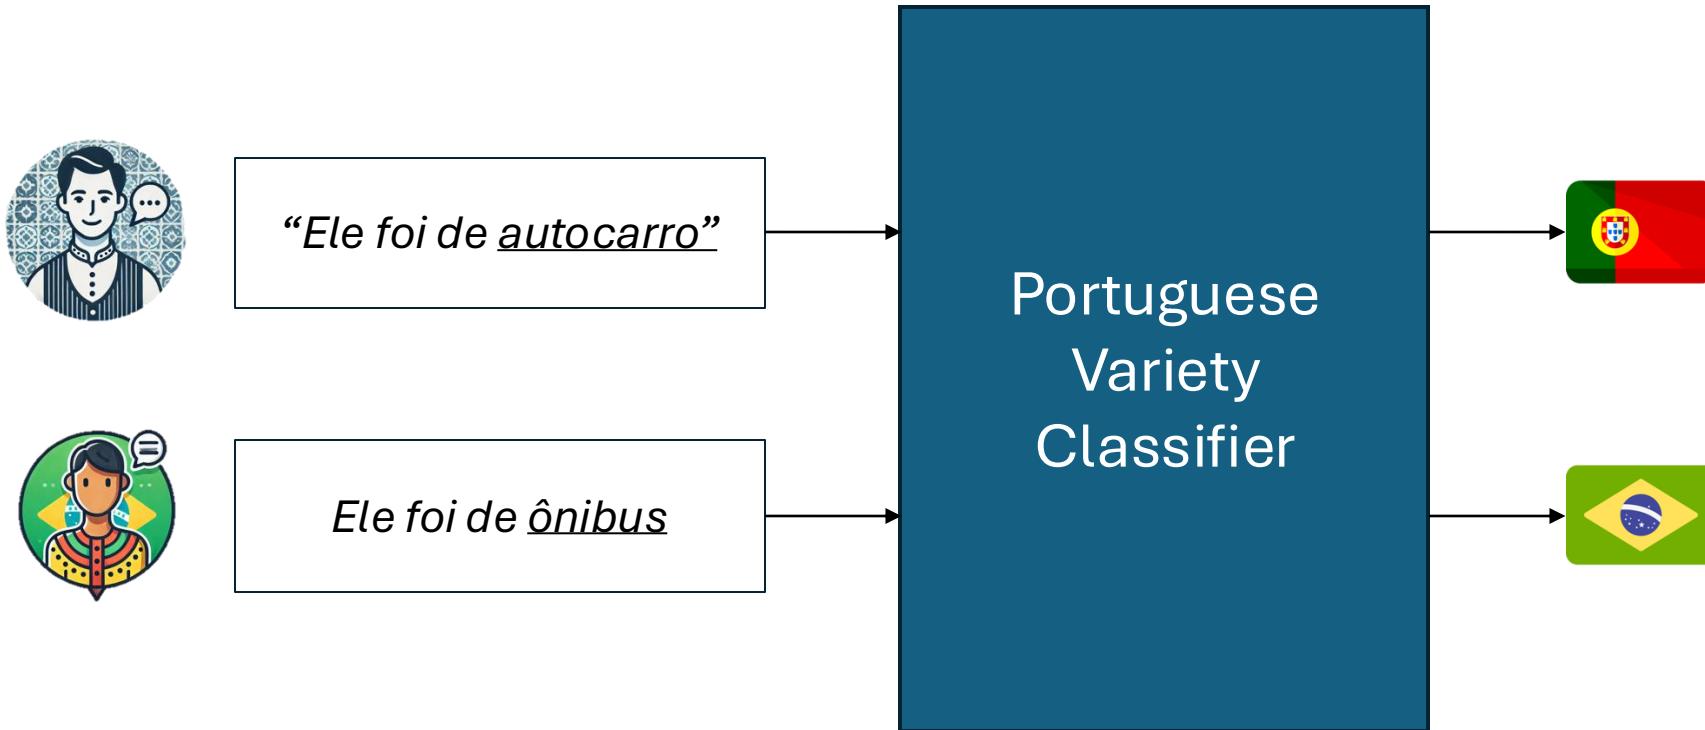
 Grok

Time is a river that looks calm from the bridge, but if you jump in you realize it's a whitewater torrent dragging you backward over rocks you never saw coming—every second a hidden rapid, every memory a bruise you only notice once you're slammed against it, and the shore you keep reaching for is always the one that's already miles behind you.

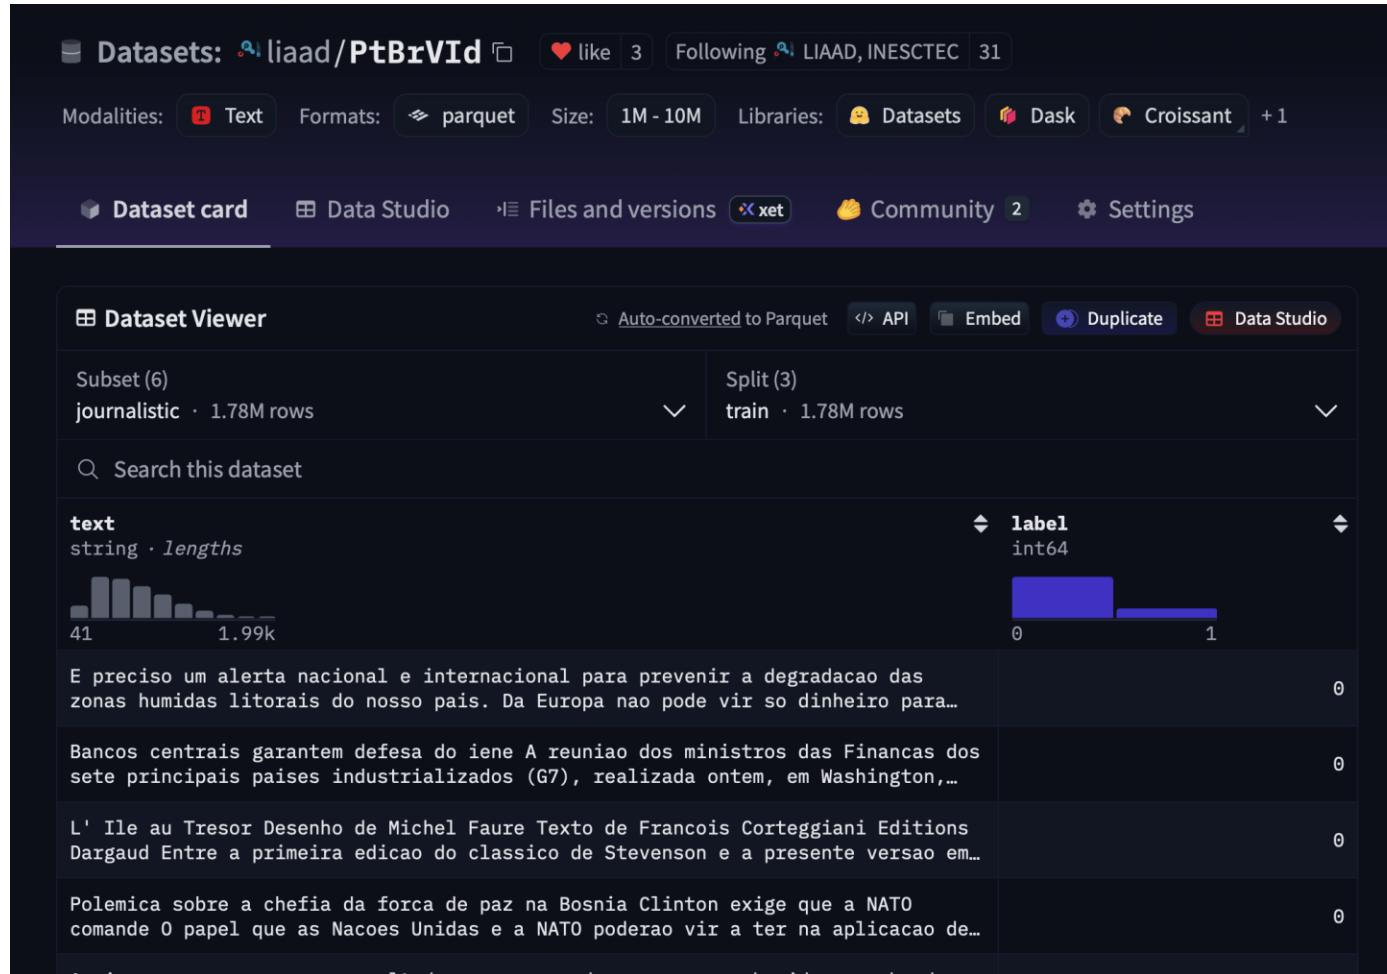
# The PtVld story

Country	Population of country <sup>[75][76][77]</sup>	More information	Native language of the majority	Spoken by
 Brazil	203,062,512	<a href="#">Portuguese in Brazil</a>	✓	95% as a native language <sup>[78]</sup>
 Angola	35,981,281	<a href="#">Portuguese in Angola</a>	✗	40% as a native language, 60% total <sup>[79]</sup>
 Mozambique	32,513,805	<a href="#">Portuguese in Mozambique</a>	✗	17% as a native language, 44% total <sup>[78]</sup>
 Portugal	10,467,366	<a href="#">Portuguese in Portugal</a>	✓	95% as a native language <sup>[80]</sup>
 Guinea-Bissau	2,078,820	<a href="#">Portuguese in Guinea-Bissau</a>	✗	0.3% as a native language, 20% total <sup>[81]</sup>

# The PtVld story: the goal



# The PtVld story: the dataset



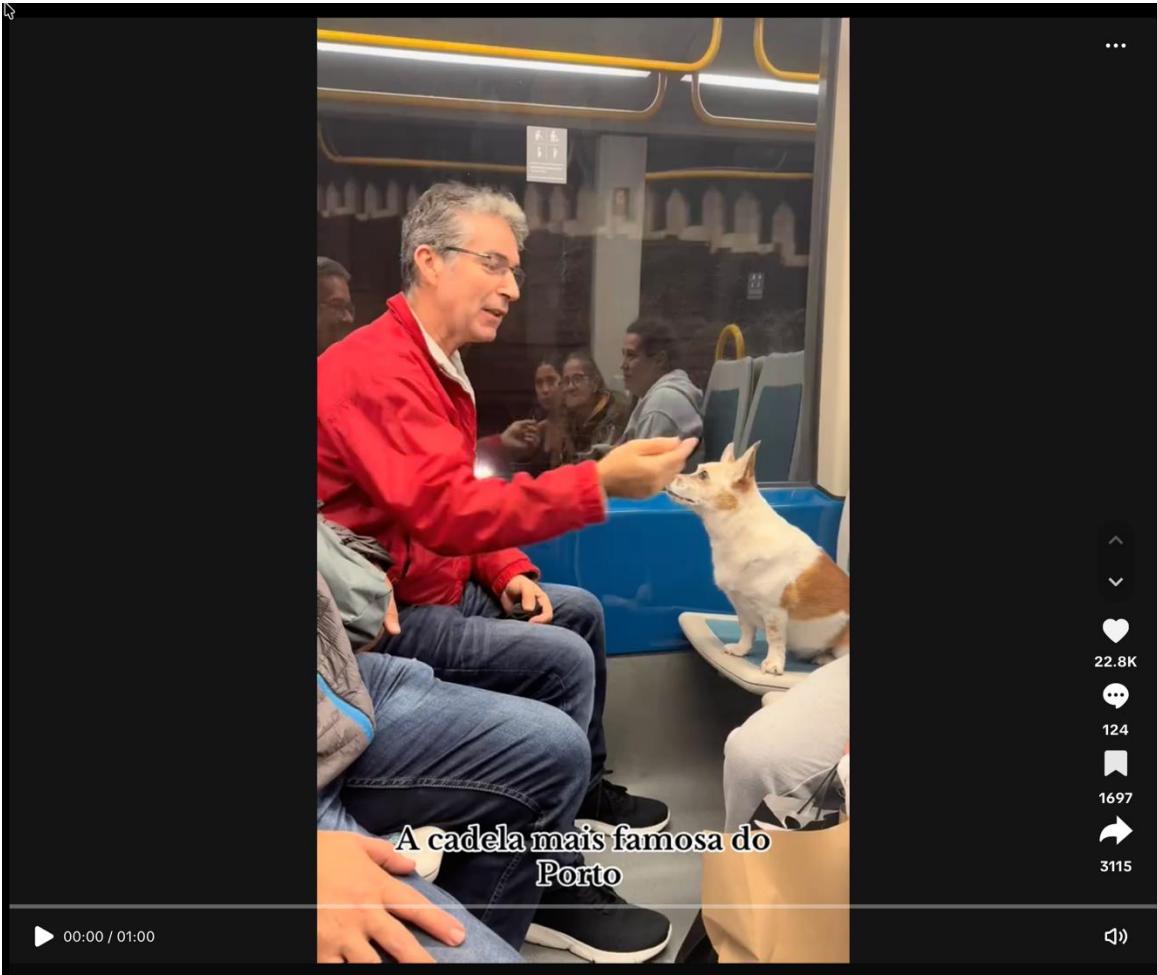
	Label	Tokens Count	Docs Count
Journalistic	PT	189,506,320	1,443,422
	BR	27,077,538	333,903
Literature	PT	1,859,660	24,090
	BR	3,805,896	52,458
Legal	PT	152,717,737	2,957,980
	BR	221,167	4,653
Politics	PT	7,203,739	27,887
	BR	1,012,586	3,656
Web	PT	22,598,587	43,630
	BR	23,913,771	44,313
Social Media	PT	44,758,304	2,363,261
	BR	94,177	5,504

<https://ojs.aaai.org/index.php/AAAI/article/view/34705>

# The PtVId story: the filtering

*Be maniacal about your code*

# The smart dog



# Clever Hans



# The same can happen to your model!

```
✓ /opt/homebrew/bin/python3  
Python 3.13.3 (main, Apr  8 2025, 13:54:08) [Clang  
Type "help", "copyright", "credits" or "license" f  
>>> hash([42])  
Traceback (most recent call last):  
  File "<python-input-0>", line 1, in <module>  
    hash([42])  
    ~~~~~^~~~~~  
TypeError: unhashable type: 'list'  
>>>
```

**TRUST ME!**

```
>>> train()  
100%|██████████| 100/100  
Training finished  
Your model is ready great!  
>>> █
```



# Training Neural Models

## Start simple

- Small architecture

## Build your training loop

- No fancy features
- No dropout, batch norm, LR decay, et cetera

## Setup Tensorboard

- Log train and valid loss of every batch
- After every train step run a forward pass on a random sample of a valid set

## Overfit to a 1 example

- The network should be able to memorize this example
- Train loss should be zero

## Overfit to a 10 example

- Now the input will be taken into consideration
- Train loss should be zero
- Val loss should go up

## Use more data

- Now valid should start to go down
- If so, the model is learning

## Profile your code

- Make your training runs go faster

optional

## Log more metrics

- ROC-curves
- Prediction distributions
- et cetera

## Hyperparameter tuning

- ideally, in parallel
- data augmentation: never apply the transformations to the valid set

## Try different approaches

- In practice, the baseline all models are trying to beat is the best approach

*Constrain as little  
as possible*

# The Bitter Lesson

Rich Sutton

March 13, 2019

<http://www.incompleteideas.net/Incldeas/BitterLesson.html>

“The biggest lesson that can be read from 70 years of AI research is that **general methods that leverage computation are ultimately the most effective**, and by a large margin.”

“Seeking an improvement that makes a difference in the shorter term, researchers **seek to leverage their human knowledge** of the domain, but **the only thing that matters in the long run is the leveraging of computation.**”





# In practice

- **Don't hard-code your own understanding of the world into the model, because your understanding is flawed and doesn't scale.**
- **Compute scales, human intuition doesn't.**
- Inductive bias:
  - Try to find ways to reduce the bias of the model so that it generalizes
  - Increasing the amount of data is usually the best way

*Just do it*

# Taste vs. Skill

- There is a gap between what we imagine we can do and what we can do

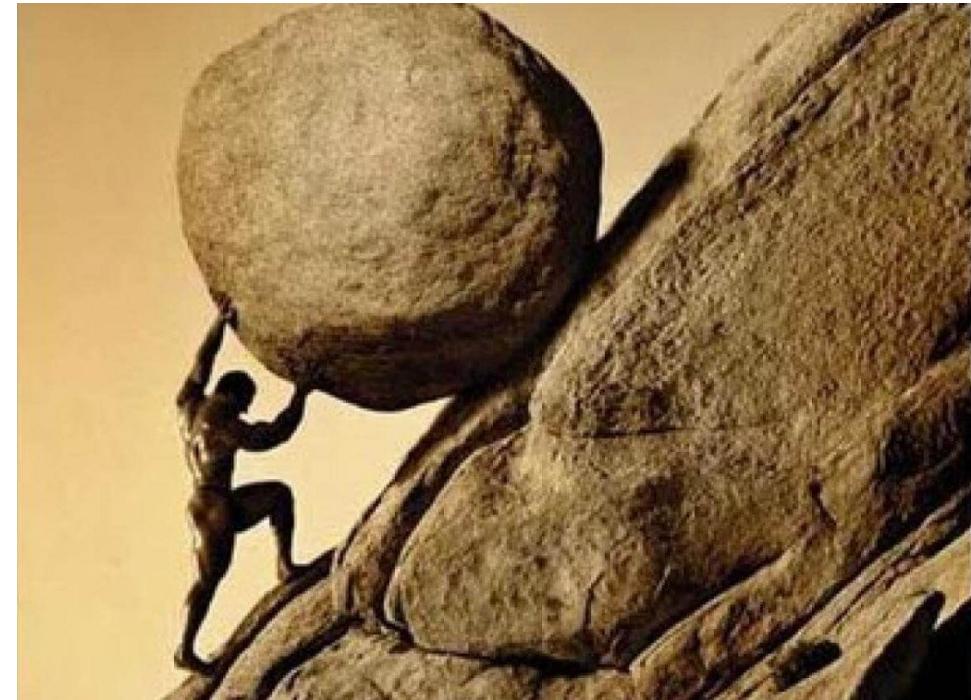
"Man is the only animal that laughs and weeps; for he is the only animal that is struck with the difference between what things are and what they ought to be." — William Hazlitt

# Procrastination

- Procrastination is a defence mechanism to protect the "perfect version" of the project in your head from becoming the "imperfect reality."
- Solution: Take action
- The photography class test

*“The best is the enemy of the good”*

# A side note: Productivity in the era of AI



*Four Thousand Weeks: Time Management for Mortals*

# A side note: Productivity in the era of AI

## The Eisenhower Decision Matrix



FOMO  
↓  
JOMO

*Be bold, be simple*

# “What is research taste?”

Ugliness, there's no room for ugliness.  
It's beauty, simplicity, elegance,

Ilya Sutskever – We're moving from the age of scaling to the age of research

Dwarkesh Patel 1.12M subscribers

Subscribe

25K | 1 Share | \$ Thanks | Clip | ...

# Some examples of simplicity

$$\vec{F} = m\vec{a}$$

$$i\hbar \frac{\partial}{\partial t} |\Psi\rangle = \hat{H} |\Psi\rangle$$

$$E=mc^2$$

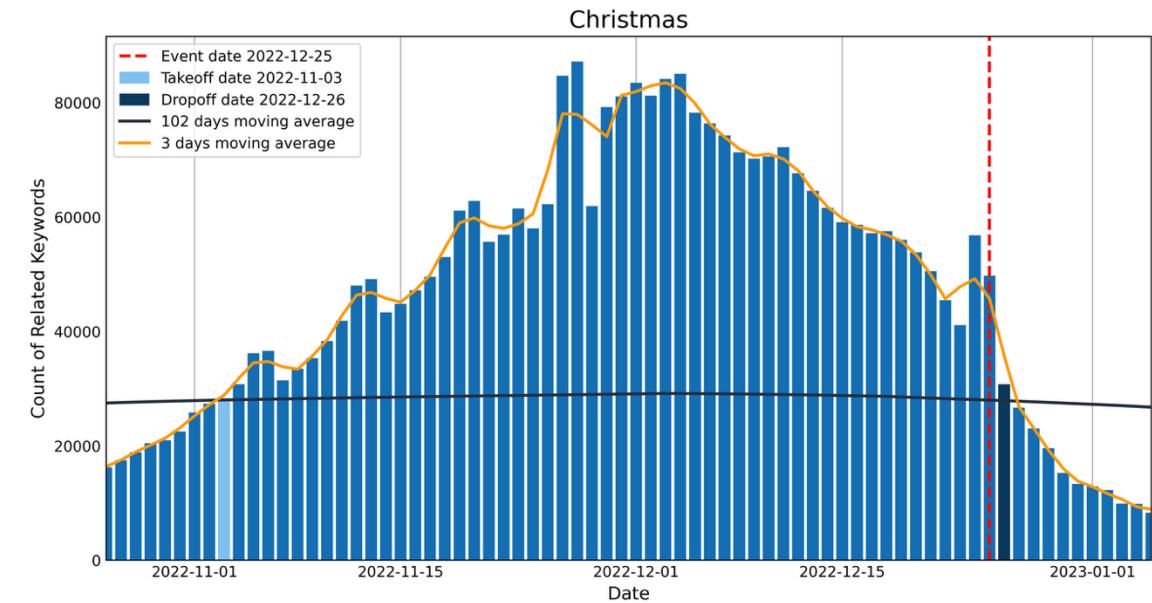
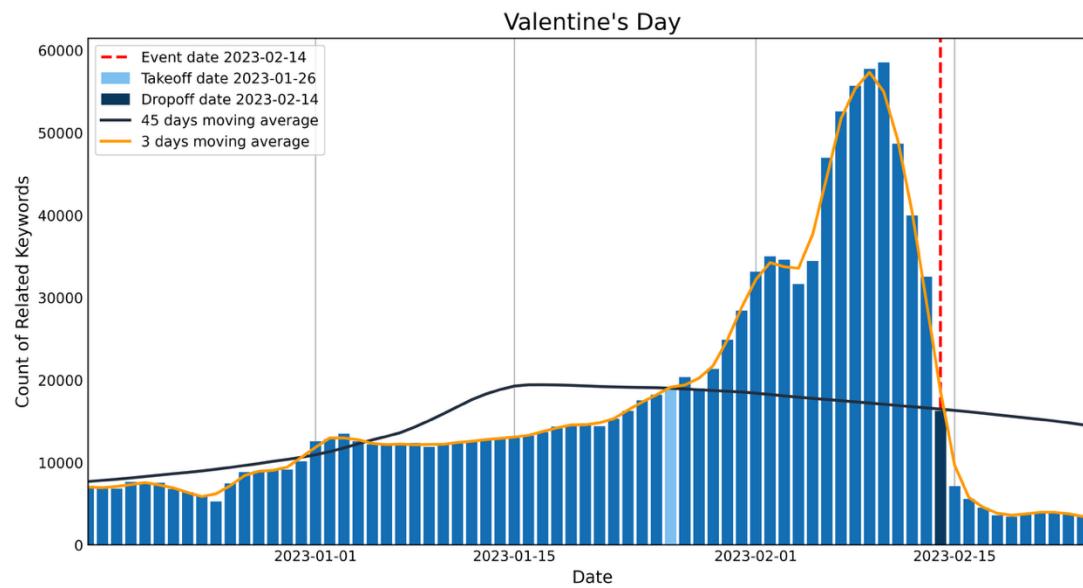
$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

$$\begin{aligned} \text{MultiHead}(Q, K, V) &= \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \\ \text{where } \text{head}_i &= \text{Attention}(QW_i^Q, KW_i^K, VW_i^V) \end{aligned}$$

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

1

# Event Temporal Models



[Don't Forget This: Augmenting Results with Event-Aware Search](#)

Thank you for you attention